

OUT-OF-SAMPLE RETURN PREDICTABILITY: A QUANTILE COMBINATION APPROACH

LUIZ RENATO LIMA^{a,b*} AND FANNING MENG^a

^a Department of Economics, University of Tennessee at Knoxville, TN, USA

^b Federal University of Paraíba, Brazil

SUMMARY

This paper develops a novel forecasting method that minimizes the effects of weak predictors and estimation errors on the accuracy of equity premium forecasts. The proposed method is based on an averaging scheme applied to quantiles conditional on predictors selected by LASSO. The resulting forecasts outperform the historical average, and other existing models, by statistically and economically meaningful margins. Copyright © 2016 John Wiley & Sons, Ltd.

Received 10 September 2015; Revised 8 August 2016



Supporting information may be found in the online version of this article.

1. INTRODUCTION

Stock return is a key variable to firms' capital structure decisions, portfolio management, asset pricing and other financial problems. As such, forecasting return has been an active research area since Dow (1920). Rapach and Zhou (2013), Campbell (2000) and Welch and Goyal (2008) illustrate a number of macroeconomic predictors and valuation ratios that are often employed in equity premium forecasting models. Valuation ratios include the dividend–price, earnings–price and book-to-market ratios, and macroeconomic variables include nominal interest rates, the inflation rate, term and default spreads, corporate issuing activity, the consumption–wealth ratio and stock market volatility. However, if such predictors are weak in the sense that their effects on the conditional mean of the equity premium are very small, then including them in the forecasting equation will result in low-accuracy forecasts that may be outperformed by the simplistic historical average (HA) model.

This paper develops a forecasting method that minimizes the negative effects of weak predictors and estimation errors on equity premium forecasts. Our approach relies on the fact that the conditional mean of a random variable can be approximated through the combination of its quantiles. This method has a long tradition in statistics and has been applied in the forecasting literature by Judge *et al.* (1988), Taylor (2007), Ma and Pohlman (2008) and Meligkotsidou *et al.* (2014). Our novel contribution to the literature is that we explore the existence of weak predictors in the quantile functions, which are identified through the ℓ_1 -penalized (LASSO) quantile regression method (Belloni and Chernozhukov, 2011). In applying such a method, we select predictors significant at the 5% level for the quantile functions. Next, we estimate quantile regressions with only the selected predictors, resulting in the post-penalized quantiles. These quantiles are then combined to obtain a point forecast of the equity premium, named the post-LASSO quantile combination (PLQC) forecast.

Our approach essentially selects a specification for the prediction equation of the equity premium. If a given predictor is useful to forecast some, but not all, quantiles of the equity premium, it is classified as partially weak. If the predictor helps forecast all quantiles, it is considered to be strong, whereas

* Correspondence to: Luiz Renato Lima, Department of Economics, 525 Stokely Management Center, University of Tennessee at Knoxville, TN, 37996-0550 USA. E-mail: llima@utk.edu

*There
are no
strong.
?*

predictors that help predict no quantile are called fully weak predictors. The ℓ_1 -penalized method sorts the predictors according to this classification. The quantile averaging results in a prediction equation in which the coefficients of fully weak predictors are set to zero, while the coefficients of partially weak predictors are adjusted to reflect the magnitude of their contribution to the equity premium forecasts. Our empirical results show that, of the 15 commonly used predictors that we examine, nine are fully weak and six are partially weak. We show that failing to account for partially weak predictors results in misspecified prediction equations and, therefore, inaccurate equity premium forecasts.

We demonstrate that the proposed PLQC method offers significant improvements in forecast accuracy not only over the historical average but also over many other forecasting models. This holds for both statistical and economic evaluations across several out-of-sample intervals. Furthermore, we develop a decomposition of the mean square prediction error (MSPE) in order to summarize the contribution of each step of the proposed PLQC approach. In other words, we measure the additional loss that would arise from weak predictors and/or the estimation errors caused by extreme observations of equity premium. In particular, our results point out that in the 1967:1–1990:12 period weak predictors explain about 15% of additional loss resultant from the non-robust forecast relative to the PLQC forecast. However, when we look at the 1991:1–2013:12 out-of-sample period, two-thirds of the loss of accuracy come from the existence of weak predictors. Not surprisingly, the forecasts that fail to account for weak predictors are exactly the ones largely outperformed by the historical average during the 1991:1–2013:12 period.

Additionally, we conduct a robustness analysis by considering quantile combination models based on known predictors.¹ These models are not designed to deal with partial and fully weak predictors across quantiles and over time. Our empirical results show that equity premium forecasts obtained by combining quantile forecasts from such models are unable to provide a satisfactory solution to the original puzzle reported by Welch and Goyal (2008).

The remainder of this paper is organized as follows. Section 2 presents the econometric methodology and introduces the quantile combination approach. It also offers a comparison of the new and existing forecasting methods. Section 3 presents the main results concerning the use of a quantile combination approach to forecast the equity premium. Section 4 concludes.

2. ECONOMETRIC METHODOLOGY

Suppose that an econometrician is interested in forecasting the equity premium² of the S&P 500 index $\{r_{t+1}\}$, given the information available at time t , I_t . The data-generating process (DGP) is defined as

$$\begin{aligned} r_{t+1} &= X'_{t+1,t} \alpha + (X'_{t+1,t} \gamma) \eta_{t+1} \\ \eta_{t+1} | I_t &\sim \text{i.i.d. } F_\eta(0, 1) \end{aligned} \quad (1)$$

where $F_\eta(0, 1)$ is some distribution with mean zero and unit variance that does not depend on I_t ; $X_{t+1,t} \in I_t$ is a $k \times 1$ vector of covariates available at time t ; $\alpha = (\alpha_0, \alpha_1, \dots, \alpha_{k-1})'$ and $\gamma = (\gamma_0, \gamma_1, \dots, \gamma_{k-1})'$ are $k \times 1$ vectors of parameters, α_0 and γ_0 being intercepts. This is the conditional location-scale model that satisfies Assumption D.2 of Patton and Timmermann (2007b) and includes most common volatility processes, e.g. autoregressive conditional heteroskedasticity and stochastic volatility. Several special cases of model (1) have been considered in the forecasting literature.³ In this

*What is
n?*

¹ We also conducted a robustness analysis using the ‘kitchen-sink’ and common-factor models. Results are shown in the online Appendix (supporting information and at <http://econ.bus.utk.edu/department/faculty/lima.asp>).

² The equity premium is calculated by subtracting the risk-free return from the return of the S&P 500 index.

³ Gaglianone and Lima (2012, 2014) assume $X_{t+1,t} = (1, C_{t+1,t})'$ and $X_{t+1,t} = (1, f_{t+1,t}^1, \dots, f_{t+1,t}^n)'$ respectively, where $C_{t+1,t}$ is the consensus forecast made at time t from the Survey of Professional Forecasts and $f_{t+1,t}^j$, $j = 1, \dots, n$, are point forecasts made at time t by different economic agents.

paper, we consider another special case of model (1) by imposing $X'_{t+1,t} = X'_t$, a vector of predictors observable at time t . In this case, the conditional mean of r_{t+1} is given by $E(r_{t+1}|X_t) = X'_t\alpha$, whereas the conditional quantile of r_{t+1} at level $\tau \in (0, 1)$, $Q_\tau(r_{t+1}|X_t)$, equals

$$Q_\tau(r_{t+1}|X_t) = X'_t\alpha + X'_t\gamma F_\eta^{-1}(\tau) = X'_t\beta(\tau) \quad (2)$$

where $\beta(\tau) = \alpha + \gamma F_\eta^{-1}(\tau)$, and $F_\eta^{-1}(\tau)$ is the unconditional quantile of η_{t+1} . Thus this model generates a linear quantile regression for r_{t+1} , where the conditional mean parameters, α , enter in the definition of the quantile parameter, $\beta(\tau)$.⁴

Following the literature (Granger 1969; Granger and Newbold 1986; Christoffersen and Diebold 1997, Patton and Timmermann 2007a), we assume that the loss function is defined as follows:

Assumption 1: loss function. The loss function L is a homogeneous function solely of the forecast error $e_{t+1} \equiv r_{t+1} - \hat{r}_{t+1}$, i.e. $L = L(e_{t+1})$ and $L(ae) = g(a)L(e)$ for some positive function g .⁵

Proposition 1 presents our result on forecast optimality. It is a special case of Proposition 3 of Patton and Timmermann (2007) in the sense that we assume a DGP with specific dynamic for the mean and variance. Under this case, we are able to show that the optimal forecast of the equity premium can be decomposed as the sum of its conditional mean and a bias measure:

Proposition 1. Under DGP (1) with $X'_{t+1,t} = X'_t$ and a homogeneous loss function (Assumption 1), the optimal forecast will be

$$\begin{aligned}\hat{r}_{t+1} &= Q_\tau(r_{t+1}|X_t) \\ &= E(r_{t+1}|X_t) + \kappa_\tau\end{aligned}$$

$$X'_t\alpha = E(r_{t+1}|X_t)$$

where $\kappa_\tau = X'_t\gamma F_\eta^{-1}(\tau)$ is a bias measure relative to the conditional mean (MSPE) forecast. This bias depends on X_t , the distribution F_η and loss function L .⁶

The above result suggests that, when estimation of the conditional mean is affected by the presence of extreme observations, as is the case with financial data, an approach to obtain robust MSPE forecasts of equity premium is through the combination of quantile forecasts. That is

$$\begin{aligned}\sum_{\tau=\tau_{\min}}^{\tau_{\max}} \omega_\tau Q_\tau(r_{t+1}|X_t) &= E(r_{t+1}|X_t) + \sum_{\tau=\tau_{\min}}^{\tau_{\max}} \omega_\tau \kappa_\tau \\ &= E(r_{t+1}|X_t) + X'_t \gamma \sum_{\tau=\tau_{\min}}^{\tau_{\max}} \omega_\tau F_\eta^{-1}(\tau)\end{aligned}$$

Biased data?
Combine quantile
forecasts instead of
simply taking
mean.

where ω_τ is the weight assigned to the conditional quantile $Q_\tau(r_{t+1}|X_t)$. Note that the weights are quantile specific since they are aimed at approximating the mean of η_{t+1} , which is zero. In the one-sample setting, integrating the quantile function over the entire domain $[0, 1]$ yields the mean of the sample distribution (Koenker 2005, p. 302). Thus, given that η_{t+1} is i.i.d., we have

⁴ Model (1) can be replaced with the assumption that the quantile function of r_{t+1} is linear. Another model that generates linear quantile regression is the random coefficient model studied by Gaglianone *et al.* (2011).

⁵ This is exactly the same Assumption L2 of Patton and Timmermann (2007). Although it rules out certain loss functions (e.g. those which also depend on the level of the predicted variable), many common loss functions are of this form, such as MSE, MAE, lin-lin and asymmetric quadratic loss.

⁶ The proof is given in the online Appendix.

$E(\eta_{t+1}) = \int_0^1 F_\eta^{-1}(t) dt = 0$.⁷ However, with a finite sample, we need to consider a grid of quantiles ($\tau_{\min}, \dots, \tau_{\max}$) and approximate $\int_0^1 F_\eta^{-1}(t) dt$ by $\sum_{\tau=\tau_{\min}}^{\tau_{\max}} \omega_\tau F_\eta^{-1}(\tau)$. The choice of the weight ω_τ reflects the potential asymmetry and excess kurtosis of the conditional distribution of η_{t+1} , F_η . In the simplest case when F_η is symmetric, assigning equal weight to quantiles located in the neighborhood of the median ($\tau = 0.5$) will suffice to guarantee that $\sum_{\tau=\tau_{\min}}^{\tau_{\max}} \omega_\tau Q_\tau(r_{t+1}|X_t) = E(r_{t+1}|X_t)$. However, when F_η is asymmetric, other weighting schemes should be used. In this paper, we consider two weighting schemes.

The robustness of this approach relies on the fact that $Q_\tau(r_{t+1}|X_t)$ are estimated using the quantile regression (QR) estimator, which is robust to estimation errors caused by occasional but extreme observations of equity premium.⁸ Since the low-end (high-end) quantiles produce a downwardly (upwardly) biased forecast of the conditional mean, another insight from the combination approach is that the point forecast $\sum_{\tau=\tau_{\min}}^{\tau_{\max}} \omega_\tau Q_\tau(r_{t+1}|X_t)$ combines oppositely biased predictions, and these biases cancel each other out. This cancelling out mitigates the problem of aggregate bias identified by Issler and Lima (2009).⁹

To our knowledge, the previous discussion is the first to provide a theoretical explanation of several empirical results that use the combination of conditional quantiles to approximate the conditional mean forecast (Judge *et al.* 1988; Taylor 2007; Ma and Pohlman 2008; Meligkotsidou *et al.* 2014). A common assumption in those papers is that the specification of the conditional quantile $Q_\tau(r_{t+1}|X_t)$ is fully known by the econometrician. However, DGP (equation (1)) is unknown. Therefore, the forecasting model based on the combination of conditional quantiles with fixed predictors is still potentially misspecified, especially when predictors are weak. In what follows, we explain how we address the problem of weak predictability in the conditional quantile function.

2.1. The ℓ_1 -Penalized Quantile Regression Estimator

Rewriting equation (2), we have the conditional quantiles of r_{t+1} :

$$Q_\tau(r_{t+1}|X_t) = \beta_0(\tau) + x_t' \beta_1(\tau) \quad \tau \in (0, 1)$$

where $\beta_0(\tau) = \alpha_0 + \gamma_0 F_\eta^{-1}(\tau)$, $\beta_1(\tau) = \alpha_1 + \gamma_1 F_\eta^{-1}(\tau)$ and x_t is a $(k - 1) \times 1$ vector of predictors (excluding the intercept).

In this paper, we identify weak predictors by employing a convex penalty to the quantile regression coefficients, leading to the ℓ_1 -penalized (LASSO) quantile regression estimator (Belloni and Chernozhukov, 2011). The LASSO quantile regression estimator solves the following problem:

$$\min_{\beta_0, \beta_1} \sum_t \rho_\tau(r_{t+1} - \beta_0(\tau) - x_t' \beta_1(\tau)) + \lambda \frac{\sqrt{\tau(1-\tau)}}{m} \|\beta_1(\tau)\|_{\ell_1} \quad (3)$$

where ρ_τ denotes the ‘tick’ or ‘check’ function defined for any scalar e as $\rho_\tau(e) \equiv [\tau - 1(e \leq 0)] e$; $1(\cdot)$ is the usual indicator function; m is the size of the estimation sample; $\|\cdot\|_{\ell_1}$ is the ℓ_1 -norm, $\|\beta_1\|_{\ell_1} = \sum_{i=1}^{k-1} |\beta_{1i}|$; $x_t = (x_{1,t}, x_{2,t}, \dots, x_{(k-1),t})'$.

LASSO first selects predictor(s) from the information set $\{x_{i,t} : i = 1, 2, \dots, (k-1)\}$ for each quantile τ at each time period t (Van de Geer 2008; Manzan 2015). As for the choice of penalty level,

⁷ Recall that $F_\eta^{-1}(\tau) = Q_\tau(\eta_{t+1})$.

⁸ The robustness of an estimator can be obtained through what is known as an influence function. Following Koenker (2005, section 2.3), the influence function of the quantile regression estimator is bounded whereas that of the ordinary least squares (OLS) estimator is not.

⁹ Aggregate bias arises when we combine predictions that are mostly upwardly (downwardly) biased. In a case like that, the averaging scheme will not minimize the forecast bias.

λ , we follow Belloni and Chernozhukov (2011) and Manzan (2015). Next, we estimate a quantile regression with the selected predictors to generate a post-LASSO quantile forecast of the equity premium in $t + 1$, denoted by $f_{t+1,t}^{\tau} = \beta_0(\tau) + x_t^{*'}\beta(\tau)$, where x_t^* is (are) the predictor(s) selected at time t by the LASSO procedure with 5% significance level. This procedure is repeated to obtain a PLQF at various $\tau \in (0, 1)$. Finally, these PLQFs are combined to obtain the post-LASSO quantile combination (PLQC) forecast, $\hat{r}_{t+1} = \sum_{j=1}^k \omega_{\tau_j} f_{t+1,t}^{\tau_j}$. The PLQC is a point (MSPE) forecast of the equity premium in $t + 1$.

2.2. An Alternative Interpretation to the PLQC Forecast

In this section, we show that the PLQC forecast can be represented by a prediction equation, which is robust to the presence of weak predictors and estimation errors.

We assume a vector of potential predictors $x_t = (1 \ x_{1,t} \ x_{2,t} \ x_{3,t})'$ available at time t and quantiles $\tau \in (\tau_1, \dots, \tau_5)$. Based on x_t and τ , we obtain PLQFs of the equity premium in $t + 1$, $f_{t+1,t}^{\tau_j}$, $j = 1, 2, \dots, 5$:

$$\begin{pmatrix} f_{t+1,t}^{\tau_1} \\ f_{t+1,t}^{\tau_2} \\ f_{t+1,t}^{\tau_3} \\ f_{t+1,t}^{\tau_4} \\ f_{t+1,t}^{\tau_5} \end{pmatrix} = \begin{pmatrix} \beta_0(\tau_1) & \beta_1(\tau_1) & 0 & 0 \\ \beta_0(\tau_2) & \beta_1(\tau_2) & 0 & 0 \\ \beta_0(\tau_3) & \beta_1(\tau_3) & 0 & 0 \\ \beta_0(\tau_4) & 0 & \beta_2(\tau_4) & 0 \\ \beta_0(\tau_5) & 0 & \beta_2(\tau_5) & 0 \end{pmatrix} \times \begin{pmatrix} 1 \\ x_{1,t} \\ x_{2,t} \\ x_{3,t} \end{pmatrix} \quad (4)$$

In this example, $x_{3,t}$ is fully weak in the population because it does not help predict any quantile. In contrast, we define $x_{1,t}$ and $x_{2,t}$ as partially weak predictors because they help predict some, but not all, quantiles. The PLQC forecast is generated based on equation (5):

$$\begin{aligned} \hat{r}_{t+1} &= \sum_{j=1}^5 \omega_{\tau_j} \beta_0(\tau_j) + \sum_{j=1}^3 \omega_{\tau_j} \beta_1(\tau_j) x_{1,t} + \sum_{j=4}^5 \omega_{\tau_j} \beta_2(\tau_j) x_{2,t} \\ &= \beta_0 + \beta_1 x_{1,t} + \beta_2 x_{2,t} \end{aligned} \quad (5) \quad \left. \right\}$$

where $\beta_0 = \sum_{j=1}^5 \omega_{\tau_j} \beta_0(\tau_j)$, $\beta_1 = \sum_{j=1}^3 \omega_{\tau_j} \beta_1(\tau_j)$ and $\beta_2 = \sum_{j=4}^5 \omega_{\tau_j} \beta_2(\tau_j)$.

Standard model selection procedures such as that proposed by Koenker and Machado (1999) are not useful for selecting weak predictors for out-of-sample forecasting. Indeed, with only three predictors, five different quantile levels, $\tau \in (\tau_1, \dots, \tau_5)$ and 300 time periods, there would potentially exist 12,000 models to be considered for estimation, which is computationally prohibitive. This is why we use the ℓ_1 -penalized quantile regression method to determine the most powerful predictors among all candidates. The ℓ_1 -penalized quantile regression method rules out the fully weak predictor $x_{3,t}$ from the prediction equation, whereas $x_{1,t}$ and $x_{2,t}$ are included but their contribution to the point forecast \hat{r}_{t+1} will reflect their partial weakness. Indeed, since the contribution of $x_{1,t}$ to predicting $f_{t+1,t}^{\tau_4}$ and $f_{t+1,t}^{\tau_5}$ is weak, our forecasting device eliminates $\beta_1(\tau_4)$ and $\beta_1(\tau_5)$ from β_1 in equation (5). The same rationale explains the absence of $\beta_2(\tau_1)$, $\beta_2(\tau_2)$ and $\beta_2(\tau_3)$ in β_2 .

Moreover, if we assume that τ_1 and τ_2 are low-end quantiles whereas τ_4 and τ_5 are high-end quantiles, the coefficient matrix (equation (4)) suggests that predictor $x_{1,t}$ is prone to make downwardly biased forecasts, whereas predictor $x_{2,t}$ is prone to make upwardly biased forecasts. These oppositely biased forecasts are then combined by equation (5) to generate a low-bias and low-MSPE point forecast. Thus we avoid the problem of aggregate bias that affects traditional forecast combination methods (Issler and Lima, 2009).

Two inefficient special cases may arise when one ignores the presence of partially weak predictors. In the first case, we estimate quantile regressions with the same predictors across $\tau \in (\tau_1, \dots, \tau_5)$ to obtain the fixed-predictor quantile regression (FQR) forecast:

$$\hat{r}_{t+1} = b_0 + b_1 x_{1,t} + b_2 x_{2,t} \quad (6)$$

where $b_0 = \sum_{j=1}^5 \omega_{\tau_j} \beta_0(\tau_j)$, $b_1 = \sum_{j=1}^5 \omega_{\tau_j} \beta_1(\tau_j)$ and $b_2 = \sum_{j=1}^5 \omega_{\tau_j} \beta_2(\tau_j)$.

The second special case corresponds to the estimation of the prediction equation (6) by OLS regression of r_{t+1} on the selected predictors, $x_{1,t}$, $x_{2,t}$, and an intercept, resulting in the fixed OLS (FOLS) forecast. Although both FQR and FOLS forecasts rule out the fully weak predictors, they do not account for the presence of partially weak predictors in the population. Moreover, the FOLS forecasts will not be robust against extreme observations, since the influence function of the OLS estimator is unbounded.

To show the relative importance of accounting for partially weak predictors and estimation errors, we consider the following decomposition:

$$\text{MSPE}_{\text{FOLS}} - \text{MSPE}_{\text{PLQC}} = [\text{MSPE}_{\text{FOLS}} - \text{MSPE}_{\text{FQR}}] + [\text{MSPE}_{\text{FQR}} - \text{MSPE}_{\text{PLQC}}] \quad (7)$$

Hence we decompose the MSPE difference between FOLS and PLQC forecasts into two elements. The first element on the right-hand side of equation (7) measures the additional loss of the FOLS forecast resulting from the OLS estimator's lack of robustness to the estimation errors, while the second element represents the extra loss caused by the presence of partially weak predictors in the population. We will apply this decomposition later in the empirical section.

2.3. Weight Selection

In this paper, we consider both time-invariant and time-variant weighting schemes. The former are simple averages of $f_{t+1,t}^\tau$. More specifically, we consider a discrete grid of quantiles $\tau \in (\tau_1, \tau_2, \dots, \tau_J)$ and set equal weights $\omega_\tau = \omega$. Two leading examples are

$$\begin{aligned} \text{PLQC}_1 &: \frac{1}{3} f_{t+1,t}^{0.3} + \frac{1}{3} f_{t+1,t}^{0.5} + \frac{1}{3} f_{t+1,t}^{0.7} \\ \text{PLQC}_2 &: \frac{1}{5} f_{t+1,t}^{0.3} + \frac{1}{5} f_{t+1,t}^{0.4} + \frac{1}{5} f_{t+1,t}^{0.5} + \frac{1}{5} f_{t+1,t}^{0.6} + \frac{1}{5} f_{t+1,t}^{0.7} \end{aligned}$$

Thus the PLQC₁ and PLQC₂ attempt to approximate the point (MSPE) forecast by assigning equal weights to a discrete set of conditional quantiles. However, the importance of quantiles in the determination of optimal forecasts may not be equal and constant over time. To address this problem, we estimate the weights from a constrained OLS regression of r_{t+1} on $f_{t+1,t}^\tau$, $\tau \in (\tau_1, \tau_2, \dots, \tau_J)$, with the following two leading examples:

$$\begin{aligned} \text{PLQC}_3 &: r_{t+1} = \sum_{\tau=\tau_1}^{\tau_3} \omega_\tau f_{t+1,t}^\tau + \varepsilon_{t+1} \quad \tau \in (0.3; 0.5; 0.7) \\ \text{s.t.} \quad \omega_{\tau_1} + \omega_{\tau_2} + \omega_{\tau_3} &= 1 \quad (8) \\ \text{PLQC}_4 &: r_{t+1} = \sum_{\tau=\tau_1}^{\tau_5} \omega_\tau f_{t+1,t}^\tau + \varepsilon_{t+1} \quad \tau \in (0.3; 0.4; 0.5; 0.6; 0.7) \\ \text{s.t.} \quad \omega_{\tau_1} + \omega_{\tau_2} + \omega_{\tau_3} + \omega_{\tau_4} + \omega_{\tau_5} &= 1 \end{aligned}$$

Similar weighting schemes have been used in the forecasting literature by Judge *et al.* (1988), Taylor (2007), Ma and Pohlman (2008) and Meligkotsidou *et al.* (2014), among others.

2.4. The Forecasting Data, Procedure and Evaluation

Before explaining the forecasting data, we introduce the standard univariate predictive regressions estimated by OLS (Welch and Goyal 2008; Rapach *et al.* 2010). They are expressed as

$$r_{t+1} = \alpha_i + \beta_i x_{i,t} + \varepsilon_{i,t+1} \quad (9)$$

where $x_{i,t}$ is a variable whose predictive ability is of interest; $\varepsilon_{i,t+1}$ is an i.i.d. error term; α_i and β_i are respectively the intercept and slope coefficients specific to model $i = 1, \dots, N$. Each univariate model i yields its own forecast of r_{t+1} labeled as $f_{t+1,t}^i = \hat{E}(r_{t+1}|X_t) = \hat{\alpha}_i + \hat{\beta}_i x_{i,t}$, where $\hat{\alpha}_i$ and $\hat{\beta}_i$ are OLS estimates of α_i and β_i .

Our data¹⁰ contain monthly observations of the equity premium to the S&P 500 index and 15 predictors, which include dividend–price ratio (DP), dividend yield (DY), earnings–price ratio (EP), dividend–payout ratio (DE), stock variance (SVAR), book-to-market ratio (BM), net equity expansion (NTIS), Treasury bill rate (TBL), long-term yield (LYT), long-term return (LTR), term spread (TMS), default yield spread (DFY), default return spread (DFR), inflation (INFL) and a moving average of earning–price ratio (E10P), from December 1926 to December 2013. Contrary to Welch and Goyal (2008), we do not lag the predictor INFL, which implies that we are assuming adaptive expectations for future price changes.¹¹

In our empirical application, we generate out-of-sample forecasts of the equity premium, r_{t+1} , using (i) 15 single-predictor regression models based on equation (9); (ii) the PLQC and FQR methods with the four weighting schemes presented above; (iii) the FOLS method; (iv) the complete subset regressions (CSR) with $k = 1, 2$ and 3. The CSR method (Elliott *et al.* (2013)) combines forecasts based on predictive regressions with k number of predictors. Hence forecasts based on CSR with $k = 1$ correspond to an equal-weighted average of all possible forecasts from univariate prediction models (Rapach *et al.*, 2010). CSR models with $k = 2$ and 3 correspond to equal-weighted averages of all possible forecasts from bivariate and trivariate prediction equations, respectively.

Following Rapach *et al.* (2010), Campbell and Thompson (2008) and Welch and Goyal (2008), among others, we use the historical average of equity premium, $\bar{r}_{t+1} = \frac{1}{J} \sum_{m=1}^J r_m$, as our benchmark model. If the information available at $X_t = (1, x_{1,t}, x_{2,t}, \dots, x_{15,t})'$ is useful to predict equity premium, the forecasting models based on X_t should outperform the benchmark.

The forecasting procedure is based on recursive estimation window (Rapach *et al.*, 2010). Our estimation window starts with 361 observations from December 1926 to December 1956 and expands periodically as we move forward. The out-of-sample forecasts range from January 1957 to December 2013, corresponding to 684 observations. In addition, forecasts that rely on time-varying weighting schemes (PLQC₃, PLQC₄, FQR₃, FQR₄) require a holdout period to estimate the weights. Thus we use the first 120 observations from the out-of-sample period as an initial holdout period, which also expands periodically. In the end, we are left with a total of 564 post-holdout out-of-sample forecasts available for evaluation.¹² In addition to the whole (long) out-of-sample period (January 1967 to December 2013), we test the robustness of our findings by considering the following out-of sample subperiods: January 1967 to December 1990, January 1991 to December 2013, and the most recent interval January 2008 to December 2013.

¹⁰ The raw data come from Amit Goyal's web page (<http://www.hec.unil.ch/agoyal/>).

¹¹ A more complete definition for each variable can be found in the online Appendix.

¹² This forecasting procedure follows exactly that adopted by Rapach *et al.* (2010)).

The first evaluation measure is the out-of-sample R^2 , R_{OS}^2 , which compares the forecast from a conditional model, \hat{r}_{t+1} , to that from the benchmark (unconditional) model \bar{r}_{t+1} (Campbell and Thompson, 2008). We report the value of R_{OS}^2 in percentage terms, $R_{OS}^2 (\%) = 100 \times R_{OS}^2$. Second, to test the null hypothesis $R_{OS}^2 \leq 0$, we apply both the Diebold and Mariano (1995) and the Clark and West (2007) tests.¹³ Lastly, to evaluate the economic value of equity premium forecasts, we calculate the *certainty equivalent return (or utility gain)*, which can be interpreted as the management fee an investor would be willing to pay to have access to the additional information provided by the conditional forecast models relative to the information available in the benchmark model.¹⁴

3. EMPIRICAL RESULTS

3.1. Out-of-Sample Forecasting Results

In Figures 1 and 2 we present time series plots of the differences between the cumulative squared prediction error for the benchmark forecast and that of each conditional forecast. This graphical analysis informs the cumulative performance of a given forecasting model compared to the benchmark model over time. When the curve in each panel increases, the conditional model outperforms the benchmark, while the opposite holds when the curve decreases. Moreover, if the curve is higher at the end of the

¹³ The Diebold and Mariano (1995) and West (1996) statistics are often used to test the null hypothesis, $R_{OS}^2 \leq 0$, among non-nested models. For nested models, such as those in this paper, Clark and McCracken (2001) and McCracken (2007) show that these statistics have a non-standard distribution. Thus the Diebold–Mariano (DM) and West tests can be severely undersized under the null hypothesis and have low power under the alternative hypothesis.

¹⁴ For more detailed information on the calculation of utility gains, refer to the online Appendix.

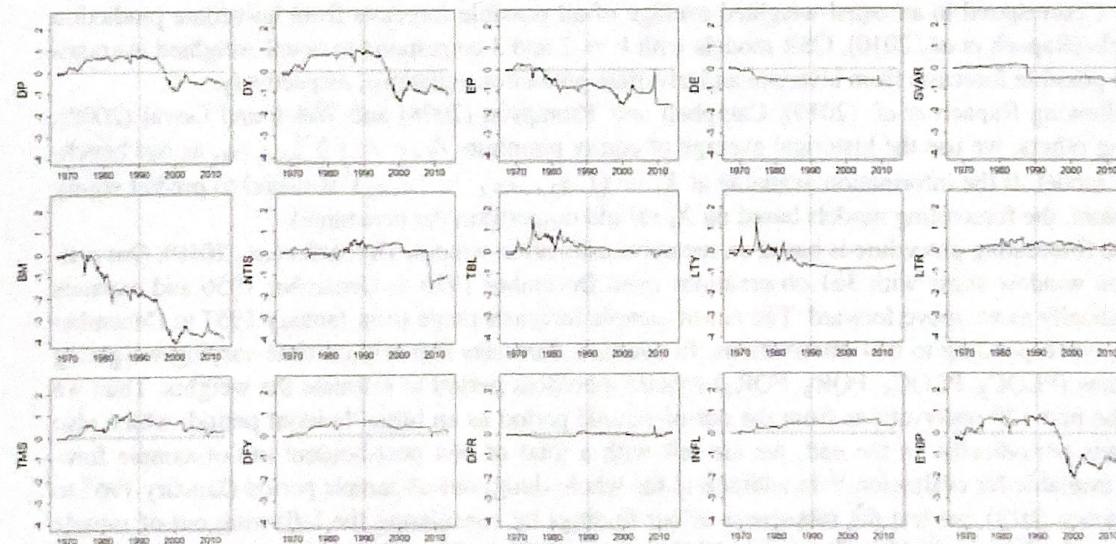
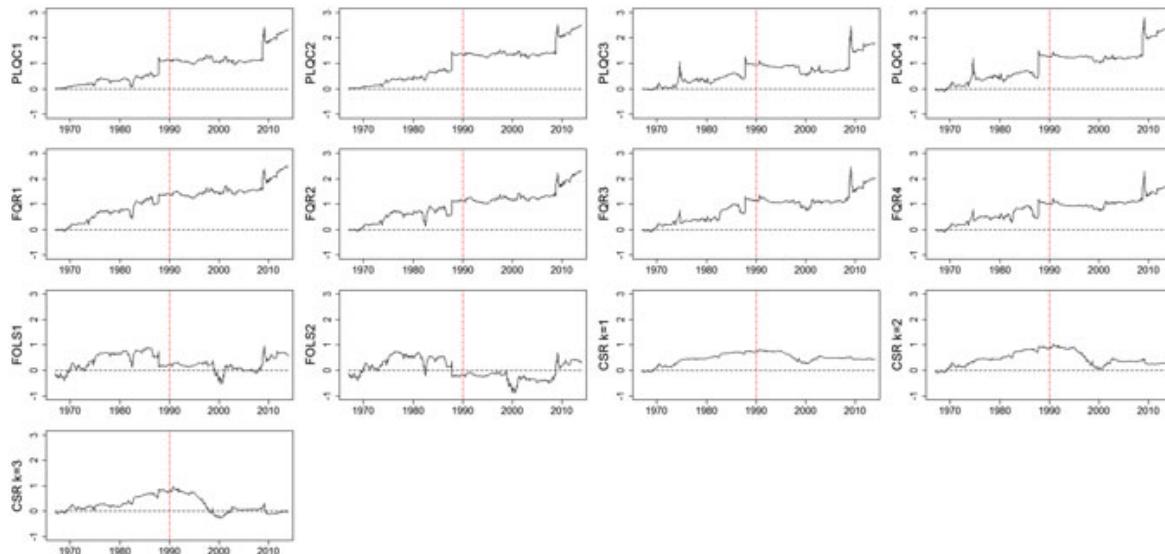


Figure 1. Cumulative squared prediction error for the benchmark model minus the cumulative squared prediction errors for the single-predictor regression forecasting models, 1967:1–2013:12. A positively sloped curve in each panel indicates that the conditional model outperforms the HA, while the opposite holds for a downward-sloping curve. Moreover, if the curve is higher at the end of the period, the conditional model has a lower MSPE than the benchmark over this period. The figure shows that, in terms of cumulative performance, few single-predictor models consistently outperform the benchmark. [Colour figure can be viewed at wileyonlinelibrary.com]

The
models
that
existed



The model proposed

Figure 2. Cumulative squared prediction error for the benchmark model minus the cumulative squared prediction errors for the FQR, FOLS, CSR and PLQC models, 1967:1–2013:12. A positively sloped curve in each panel indicates that the conditional model outperforms the HA, while the opposite holds for a downward-sloping curve. Moreover, if the curve is higher at the end of the period, the conditional model has a lower MSPE than the benchmark over this period. Figure 2 shows that the PLQC forecasts are among the top performers, especially after 1990. [Colour figure can be viewed at [wileyonlinelibrary.com](#)]

period, the conditional model has a lower MSPE than the benchmark over the whole out-of-sample period.

In general, Figure 1 shows that in terms of cumulative performance few single-predictor models consistently outperform the historical average.¹⁵ A number of the panels (such as the one based on TMS) exhibit increasing predictability in the first half of the sample period, but lose predictive strength thereafter. Also, the majority of the single-predictor forecasting models have a higher MSPE than the benchmark. Figure 1 looks very similar to that in Rapach *et al.* (2010, p. 833), which uses quarterly data. Our results, which are based on monthly observations, show a significant deterioration of the single-predictor models after 1990.¹⁶ In sum, Figure 1 strengthens the arguments already reported throughout the literature (Welch and Goyal 2008; Rapach *et al.* 2010), that it is difficult to identify individual predictors that help improve equity premium forecasts over time.

Figure 2 shows the same graphical analysis for PLQC_j , FQR_j , $j = 1, 2, 3, 4$, FOLS_1 , FOLS_2 ¹⁷ and CSR with $k = 1, 2, 3$. The curves for PLQC_j and FQR_j do not exhibit substantial falloffs like those observed in the single-predictor forecasting models (equation (9)). This indicates that the PLQC_j and FQR_j forecasts deliver out-of-sample gains on a considerably more consistent basis over time. The PLQC and FQR forecasts perform similarly, and FQR forecasts are only slightly better before 1990. Since the PLQC method accounts for partially weak predictors whereas FQR does not (this being the only difference between the two), the results shown in Figure 2 suggest that most of the predictors are

¹⁵ One exception is the single-predictor model based on INFL. Its curve is sloped upward for most of the time.

¹⁶ Welch and Goyal (2008), as well as Rapach *et al.* (2010), considered quarterly forecasts of the equity premium.

¹⁷ Recall that FOLS forecasts are based on the OLS estimation of an equation whose predictors are selected by the ℓ_1 -penalized quantile regression method. Since we have considered two sets of quantiles $\tau = (0.3, 0.5, 0.7)$ and $\tau = (0.3, 0.4, 0.5, 0.6, 0.7)$, there will be two such prediction equations and therefore two FOLS forecasts, denoted by FOLS_j , $j = 1, 2$.

not weak until 1990. The results in Figure 2 also provide the first empirical evidence concerning the ability of the PLQC model to efficiently predict the monthly equity premium of the S&P 500 index.¹⁸

The comparison between FQR and FOLS shows the importance of using quantile regression to obtain a robust estimation of the prediction equation. Recall that FQR and FOLS rely on the same specification for the prediction equation, but they differ in how the coefficients are estimated. Comparing the panels corresponding to FQR and FOLS forecasts, we see how estimation errors in the prediction equation can result in a severe loss of forecasting accuracy. The curves of the FOLS forecasts are not only lower in magnitude but also much more erratic than those corresponding to the FQR forecasts. Finally, the CSR forecasts do not outperform the PLQC forecast. Besides being robust to the presence of weak predictors and estimation errors, the PLQC forecast results from the combination of different quantile forecasts, whose biases cancel out each other. This avoids the aggregate bias problem that affects most existing forecast combination methods, including the CSR model (Issler and Lima, 2009).

We next turn to the analysis of all four out-of-sample periods. The results are displayed in Table I. This table reports R^2_{OS} statistics and its significance through the p -values of the Clark and West (2007) test (CW). It also displays the annual utility gain Δ (annual%) associated with each forecasting model and the p -value of the DM test. The results for the entire 1967:1–2013:12 out-of-sample period confirm that few single-predictor forecasting models have positive and significant R^2_{OS} . The same thing happens to the CSR forecasts. The only exceptions in this long out-of-sample period are the PLQC and FQR forecasts. Their performance is similar in the sense that they both outperform the FOLS forecast in terms of R^2_{OS} and utility gain Δ (annual%). Among the PLQC forecasts, we note that those that rely on the combination of five quantiles perform better than those based on the combination of three quantiles during this period.

As for the subperiod 1967:1–1990:12, Table I shows that some single-predictor models perform well. In particular, forecasts from single-predictor models using either DY or E10P present positive and significant R^2_{OS} and also sizable utility gains. The CSR forecasts are also reasonable and outperform the FOLS forecast. Recall that the difference between PLQC and FQR is that the latter ignores partially weak predictors, and therefore the result reported by Table I suggests that there is no advantage in using a forecasting device that is robust to (partially) weak predictors when predictors are actually strong. However, since FQR outperforms OLS-based FOLS, we conclude that forecasts which are robust against estimation errors provide a predictive advantage.

As for the 1991:1–2013:12 subperiod, we note that the R^2_{OS} of all single-predictor models fall substantially and become non-significant, suggesting that most of the predictors become weak after 1990. The same results for CSR forecasts indicate that this methodology is also affected by the presence of weak predictors. On the other hand, the results in Table I show that the R^2_{OS} of the PLQC forecast does not fall much across the two subperiods, confirming that this method is robust to weak predictors. Also, the R^2_{OS} of the FQR forecasts falls on average by 0.22%, whereas the R^2_{OS} of the PLQC forecasts increases on average by 0.18%. This happens because the latter is robust to both fully and partially weak predictors, whereas the former is only robust to fully weak predictors.

Finally, we look at the most recent out-of-sample subperiod, 2008:1–2013:12, characterized by the occurrence of the subprime crisis in the USA. A practitioner should expect that a good forecasting model would work reasonably well in periods of financial turmoil. However, the results in Table I suggest that none of the single-predictor models or the CSR forecasts perform well during this period of financial instability. In contrast, the statistic and economic measures of the PLQC forecasts are even better than those in other periods. More specifically, the R^2_{OS} and utility gain statistics for $PLQC_j$ are at least twice as large as those for other out-of-sample periods. This suggests that the PLQC method works very well even during periods with multiple episodes of financial turmoil. These results provide

¹⁸ Based on a Monte Carlo simulation experiment, we found that weak predictors can be harmful for forecasting. Our Monte Carlo experiment can be found in the online Appendix.

Table I. Out-of-sample equity premium forecasting.

Model	OOS: 1967:1–2013:12			OOS: 1967:1–1990:12			OOS: 1991:1–2013:12			OOS: 2008:1–2013:12		
	$R_{OS}^2(\%)$	DM	CW	$\Delta(\text{annual}\%)$	$R_{OS}^2(\%)$	DM	CW	$\Delta(\text{annual}\%)$	$R_{OS}^2(\%)$	DM	CW	$\Delta(\text{annual}\%)$
<i>Single-predictor model forecasts</i>												
DP	-0.60	1.00	0.26	-0.10	1.31	0.30	0.03	1.72	-2.99	1.00	0.78	-1.99
DY	-1.01	1.00	0.22	0.22	1.55	0.38	0.02	2.16	-4.24	1.00	0.76	-1.79
EP	-1.51	1.00	0.48	-0.29	-0.99	0.93	0.43	-0.67	-2.15	0.99	0.53	0.11
DE	-0.71	0.98	0.98	-0.55	-0.90	1.00	1.00	-0.95	-0.48	0.73	-0.14	-1.15
SVAR	-0.54	0.71	0.81	-0.26	-0.74	0.57	0.10	-0.29	0.94	0.83	-0.43	-0.77
BM	-3.54	1.00	0.62	-1.43	-2.75	0.95	0.44	-0.97	-4.53	1.00	0.79	-1.90
NTIS	-0.96	0.58	0.36	-1.12	0.38	0.54	0.09	-0.34	-2.65	0.58	0.83	-1.92
TBL	0.09	0.24	0.07	2.09	0.37	0.27	0.06	4.11	-0.26	0.29	0.53	-0.01
LTY	-0.64	0.40	0.14	1.84	-1.08	0.38	0.13	3.60	-0.09	0.65	0.54	0.01
LTR	0.12	0.81	0.12	0.25	0.55	0.69	0.09	1.16	-0.43	0.79	0.42	-0.71
TMS	0.28	0.42	0.07	0.86	1.26	0.43	0.03	1.90	-0.96	0.46	0.53	-0.23
DFY	0.13	0.73	0.22	0.01	1.00	0.10	0.01	1.14	-0.97	0.99	0.87	-1.17
DFR	0.04	0.55	0.36	0.06	0.01	0.60	0.43	0.10	0.08	0.52	0.37	0.02
INFL	0.37	0.01	0.10	0.69	0.78	0.06	0.07	1.47	-0.14	0.03	0.51	-0.13
E10P	-1.42	1.00	0.17	0.06	1.32	0.58	0.04	1.84	-4.86	1.00	0.66	-1.78
<i>Complete subset regression forecasts</i>												
CSR $k = 1$	0.39	0.86	0.06	0.49	1.29	0.07	0.00	1.58	-0.75	1.00	0.82	-0.65
CSR $k = 2$	0.24	0.96	0.11	0.46	1.61	0.23	0.00	1.96	-1.48	1.00	0.83	-1.11
CSR $k = 3$	-0.02	0.99	0.20	0.39	1.49	0.45	0.02	1.84	-1.93	1.00	0.82	-1.12
<i>Forecasts based on LASSO-quartile selection</i>												
FOLS1	0.53	0.50	0.03	1.35	0.45	0.58	0.11	1.12	0.63	0.43	0.09	1.58
FOLS2	0.27	0.62	0.04	1.18	-0.19	0.75	0.17	0.71	0.84	0.40	0.07	1.67
FQR1	2.27	0.00	0.00	2.42	2.34	0.15	0.01	2.23	2.20	0.00	0.02	2.60
FQR2	2.10	0.02	0.00	2.25	1.96	0.38	0.02	2.15	2.29	0.01	0.02	2.34
FQR3	1.83	0.07	0.01	2.07	2.04	0.18	0.04	2.46	1.56	0.13	0.08	1.65
FQR4	1.56	0.13	0.02	1.83	1.82	0.33	0.05	2.27	1.24	0.12	0.11	1.37
PLQCI	2.12	0.01	0.01	1.59	1.81	0.16	0.05	1.03	2.50	0.02	0.04	2.19
PLQC2	2.27	0.00	0.01	1.86	2.23	0.09	0.04	1.76	2.31	0.01	0.04	1.96
PLQC3	1.62	0.09	0.04	1.69	1.61	0.13	0.10	2.46	1.62	0.21	0.11	1.79
PLQC4	2.16	0.06	0.03	2.16	2.20	0.16	0.08	2.97	2.11	0.11	0.08	1.32

Note: This table reports R_{OS}^2 statistics (in %) and its significance through the p -values of the Clark and West (2007) test (CW). It also reports the p -value of the Diebold and Mariano (1995) test (DM) and the annual utility gain $\Delta(\text{annual}\%)$ associated with each forecasting model over four out-of-sample periods. $R_{OS}^2 > 0$, if the conditional forecast outperforms the benchmark. The annual utility gain is interpreted as the annual management fee that an investor would be willing to pay in order to get access to the additional information from the conditional forecast model.

strong evidence that we have identified an effective method for forecasting monthly equity premium on the S&P 500 index based on economic variables.

Table II shows the decomposition of the MSPE introduced in Section 2.2. Recall that this decomposition measures the additional MSPE loss of FOLS forecasts relative to the PLQC forecasts. The first element on the right-hand side of equation measures the additional loss of the FOLS forecast resulting from the OLS estimator's lack of robustness to the estimation errors, while the second element represents the extra loss caused by the presence of partially weak predictors in the population. For the 1967:1–1990:12 subperiod, the contribution of partially weak predictors is much smaller compared to that of estimation errors. This is consistent with the results shown in Figures 1 and 2 and also those in Table I. In the case of strong predictors, most of the loss will be explained by the OLS estimator's lack of robustness to estimation errors, so using quantile regression presents an advantage in that it avoids the effect of estimation errors. The situation changes dramatically when weak predictors become a more severe issue during the post-1990 out-of-sample period. As a result, the second element dominates, indicating that most of the forecast accuracy loss is ascribed to the presence of partially weak predictors.

In the next section we provide more information that explains the benefits of the PLQC forecasts.

Table II. Mean squared prediction error (MSPE) decomposition

$\text{MSPE}_{\text{FOLS}} - \text{MSPE}_{\text{PLQC}}$	$(\text{MSPE}_{\text{FOLS}} - \text{MSPE}_{\text{FQR}}) +$	$(\text{MSPE}_{\text{FQR}} - \text{MSPE}_{\text{PLQC}})$
OOS	% of total	% of total
1967:1–1990:12	84.3%	15.7%
1991:1–2013:12	31.3%	68.7%

Note: The decomposition measures the additional MSPE loss of FOLS forecasts relative to the PLQC forecasts. The first element ($\text{MSPE}_{\text{FOLS}} - \text{MSPE}_{\text{FQR}}$) measures the additional loss from the OLS estimator's lack of robustness to estimation errors, while the second element ($\text{MSPE}_{\text{FQR}} - \text{MSPE}_{\text{PLQC}}$) represents the extra loss caused by the presence of partially weak predictors in the population. Note that the PLQC, FQR and FOLS forecasts correspond to models noted as PLQC4, FQR4 and FOLS2 in the paper.

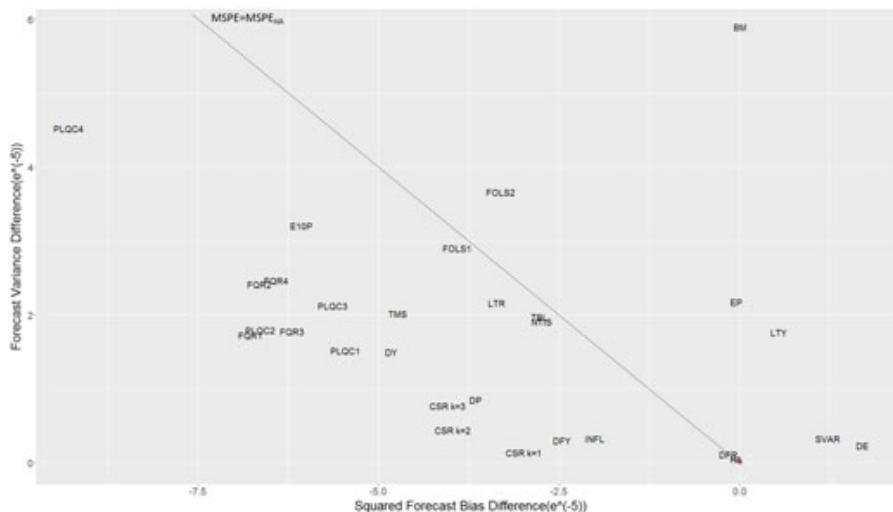


Figure 3. Scatter plot of forecast variance and squared forecast bias relative to historical average, 1967:1–1990:12. The y -axis and x -axis represent relative forecast variance and squared forecast bias of all single-predictor models, CSR, FOLS, FQR and PLQC models, calculated as the difference between the forecast variance (squared bias) of the conditional model and the forecast variance (squared bias) of the HA. Each point on the dotted line represents a forecast with the same MSPE as the HA; points to the right are forecasts outperformed by the HA, and points to the left represent forecasts that outperform the HA

3.2. Explaining the Benefits of the PLQC Forecasts

In this section we decompose the MSPE into two parts: the forecast variance and the squared forecast bias. We calculate the MSPE of any forecast \hat{r}_{t+1} as $\frac{1}{T^*} \sum_t (r_{t+1} - \hat{r}_{t+1})^2$ and the unconditional forecast variance as $\frac{1}{T^*} \sum_t (\hat{r}_{t+1} - \frac{1}{T^*} \sum_t \hat{r}_{t+1})^2$, where T^* is the total number of out-of-sample forecasts. The squared forecast bias is computed as the difference between MSPE and forecast variance (Elliott *et al.*, 2013; Rapach *et al.*, 2010).

Figures 3 and 4 depict the relative forecast variance and squared forecast bias of all single-predictor models, CSR, FOLS, FQR and PLQC models for two out-of-sample subperiods: 1967:1–1990:12 and 1991:1–2013:12. The relative forecast variance (squared bias) is calculated as the difference between the forecast variance (squared bias) of the i th model and the forecast variance (squared bias) of the historical average (HA). Hence the value of relative forecast variance (squared bias) for the HA is necessarily equal to zero. Each point on the dotted line represents a forecast with the same MSPE as the HA; points to the right of the line are forecasts outperformed by the HA, and points to the left represent forecasts that outperform the HA. Finally, both forecast variance and squared forecast bias are measured at the same scale, so it is possible to determine the trade-off between variance and bias of each forecasting model.

Since the HA forecast is a simple average of historical equity premium, it will have a very low variance but will be biased. Figure 3 shows that, in the 1967:1–1990:12 subperiod, most of the forecasts based on single-predictor models outperformed the HA. Combining this result with the empirical observation that the variances of forecasts based on single-predictor models are not lower than the variance of the HA, we conclude that such performance relies almost exclusively on a predictor's ability to lower forecast bias relative to that of HA. As a result, a predictor is classified as exhibiting strong predictability if it can produce forecasts in which the reduction in bias is greater than the increase in variance, relative to the HA forecast.

The preceding discussion offers an explanation of the results presented in Figure 4. For the sub-period 1991:1–2013:12, almost all single-predictor models are outperformed by the HA, suggesting

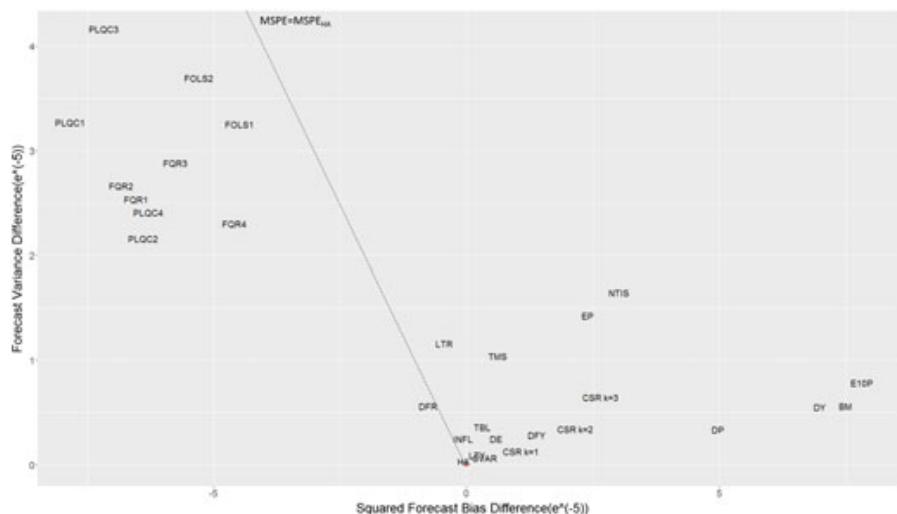


Figure 4. Scatter plot of forecast variance and squared forecast bias relative to historical average, 1991:1–2013:12. The y -axis and x -axis represent relative forecast variance and squared forecast bias of all single-predictor models, CSR, FOLS, FQR and PLQC models, calculated as the difference between the forecast variance (squared bias) of the conditional model and the forecast variance (squared bias) of the HA. Each point on the dotted line represents a forecast with the same MSPE as the HA; points to the right are forecasts outperformed by the HA, and points to the left represent forecasts that outperform the HA

the presence of weak predictors. This weak performance is mainly driven by the substantial increase in the squared biases of such forecasts. We note that when predictors are strong (in Figure 3), PLQC and FQR perform equally well. However, when predictors become weak (in Figure 4), the PLQC outperforms other forecasting methods.

Overall, the success of the PLQC forecast is explained by its ability to substantially reduce the squared forecast bias at the expense of a moderate increase in forecast variance. Additional reduction in the forecast variance of the PLQC forecasts can be obtained by increasing the number of quantiles used in the combination, as shown by points PLQC₂ and PLQC₄ in Figures 3 and 4. The main message is that the forecasting models that yield a sizable reduction in the forecast bias while keeping variance under control are able to improve forecast accuracy over HA. This explains the superior performance of PLQC forecasts.

Another analysis that we find interesting is the identification of which predictors are chosen by the ℓ_1 -penalized method across quantiles and over time. This analysis was originally suggested by Pesaran and Timmermann (1995) for the mean function. Table III shows the frequency with which each predictor is selected over the out-of-sample period, 1967:1–2013:12, and across the quantiles used to compute the PLQC forecast, i.e. $\tau = 0.3, 0.4, 0.5, 0.6$ and 0.7 . Recall from Section 2 that a predictor is defined to be partially weak if it is useful to forecast some, but not all, quantiles of the equity premium. If it helps forecast all quantiles, it is considered to be strong, whereas if it helps predict no quantile, it is fully weak. Note that Table III reports selection frequency for only six predictors, meaning that nine (out of 15) predictors are fully weak. Thus the prediction equation (5) that results from this selection procedure will include at most six predictors, but these predictors are not equally important due to their different levels of partial weakness. For instance, the selection frequency for DFY is no more than 1% at some quantiles, whereas the predictor INFL seems to be strong at almost all quantiles, except $\tau = 0.7$. Failing to account for partially weak predictors results in misspecified prediction equations and, therefore, inaccurate forecasts of equity premium as shown before.

Figure 5 shows in detail how the proposed selection procedure works over time and across quantiles. There are five charts: one for each quantile used to compute the PLQC forecast. For each chart, we list 15 predictors on the vertical axis. The horizontal axis shows the out-of-sample period. Dots inside the charts indicate that a predictor was selected to forecast a given quantile of the equity premium at time t . Figure 5 shows that predictor INFL is useful for forecasting almost all quantiles until 2010 (with noted exceptions at $\tau = 0.7$), but it loses predictability power after that. Other predictors, such as LTR, BM and SVAR, are not important at the beginning of the period but become useful for forecasting after 1985, whereas predictor E10P seems to be very useful only for forecasting the two most extreme quantiles $\tau = 0.3$ and $\tau = 0.7$.

Table III. Frequency of variables selected over OOS: January 1967 to December 2013

$\tau =$	SVAR	BM	LTR	DFY	INFL	E10P
30th	63.30%	0.18%	—	0.18%	79.61%	11.35%
40th	81.56%	—	—	—	100.00%	—
50th	40.78%	3.37%	—	—	88.65%	0.89%
60th	—	—	32.98%	—	89.18%	—
70th	—	20.74%	34.04%	—	—	55.32%

Note: The table presents the frequency with which each predictor is selected over the out-of-sample period 1967:1–2013:12 and across quantiles ($\tau = 0.3, 0.4, 0.5, 0.6$, and 0.7).

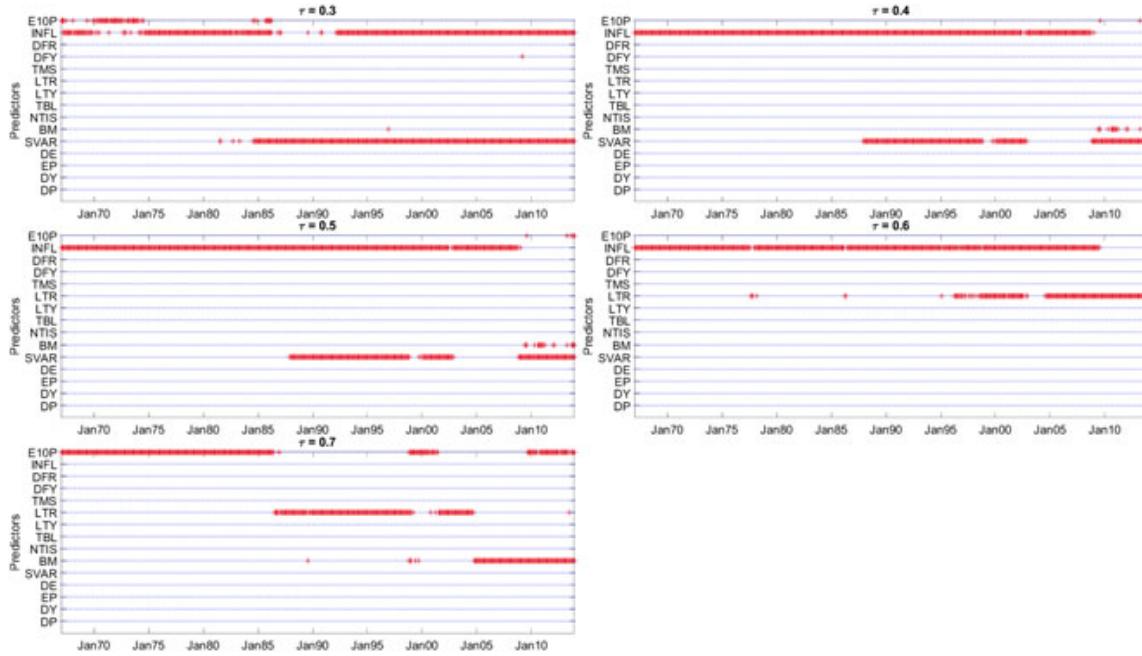


Figure 5. Variables selected by PLQC for quantile levels $\tau = 0.3, 0.4, 0.5, 0.6, 0.7$ over OOS 1967:1–2013:12. The five charts—one for each quantile used in the PLQC forecast—display the selected predictor(s) at each time point t over the out-of-sample period, 1967:1–2013:12. [Colour figure can be viewed at wileyonlinelibrary.com]

Thus, by carefully excluding fully weak predictors and identifying the relative importance of partially weak predictors, our forecasting approach can yield much better out-of-sample forecasts, which helps us understand why models that overlook weak predictors are outperformed by the proposed PLQC method.

3.3. Robustness Analysis: Other Quantile Forecasting Models

Meligkotsidou *et al.* (2014) propose the asymmetric-loss LASSO (AL-LASSO) model, which estimates the conditional quantile function as a weighted sum of quantiles by using LASSO to select the weights, i.e.

$$\theta_t = \arg \min_{\theta_t \in R^{15}} \sum_t \rho_\tau \left(r_{t+1} - \sum_{i=1}^{15} \theta_{i,t} \hat{r}_{i,t+1}(\tau) \right) \text{ s.t. } \sum_{i=1}^{15} \theta_{i,t} = 1; \quad \sum_{i=1}^{15} |\theta_{i,t}| \leq \delta_1 \quad (10)$$

Another LASSO-based quantile forecasting

where $\rho_\tau(\cdot)$ is the asymmetric loss, $\hat{r}_{i,t+1}(\tau)$ is the quantile function obtained from a single-predictor quantile model, i.e. $\hat{r}_{i,t+1}(\tau) = \alpha_i(\tau) + \beta_i(\tau)x_{i,t}$ and $x_{i,t} \in X_t = (x_{1,t}, \dots, x_{15,t})'$. The parameter δ_1 controls for the level of shrinkage. A solution to problem (10) results in an estimation of the τ th conditional quantile of r_{t+1} , $\hat{r}_{t+1}(\tau) = \sum_{i=1}^{15} \hat{\theta}_{i,t} \hat{r}_{i,t+1}(\tau)$. This process is repeated for every $\tau \in (0.3, 0.4, 0.5, 0.6, 0.7)$.

As the first robustness test, we investigate whether our PLQF, $f_{t+1,t}^\tau$, outperform other single-predictor quantile forecasts $\hat{r}_{i,t+1}(\tau)$, $i = 1, \dots, 15$ and the AL-LASSO based on the quantile score (QS) function (Manzan, 2015). The QS represents a local out-of-sample evaluation of the forecasts in the sense that, rather than providing an overall assessment of the distribution, it concentrates on a specific quantile. The higher the QS, the better the model does in forecasting a given quantile. It is computed as

Quantile Score (QS)
 is the local out-of-sample
 score at a given
 quantile.

Table IV. Quantile scores

No. model	QS ($\times 10^{-2}$) across quantile levels τ				
	$\tau =$	0.3	0.4	0.5	0.6
PLQF	-1.499	-1.641	-1.672	-1.615	-1.457
AL-LASSO	-1.532	-1.652	-1.708	-1.637	-1.464
DP	-1.532	-1.675	-1.692	-1.626	-1.459
DY	-1.533	-1.676	-1.692	-1.624	-1.460
EP	-1.539	-1.677	-1.695	-1.626	-1.461
DE	-1.568	-1.697	-1.698	-1.625	-1.450
SVAR	-1.512	-1.660	-1.688	-1.628	-1.449
BM	-1.526	-1.674	-1.694	-1.632	-1.467
NTIS	-1.527	-1.672	-1.689	-1.623	-1.449
TBL	-1.527	-1.658	-1.679	-1.619	-1.462
LTY	-1.529	-1.664	-1.689	-1.625	-1.467
LTR	-1.536	-1.678	-1.696	-1.617	-1.448
TMS	-1.532	-1.668	-1.689	-1.626	-1.462
DFY	-1.537	-1.673	-1.690	-1.618	-1.446
DFR	-1.523	-1.670	-1.692	-1.621	-1.454
INFL	-1.517	-1.648	-1.674	-1.610	-1.445
E10P	-1.529	-1.673	-1.696	-1.627	-1.466

Note: The table shows the QS for each single-predictor quantile model, AL-LASSO and PLQF models. Quantile scores are always negative. Thus the larger the QS is, i.e. the closer it is to zero, the better. The quantile scores of AL-LASSO are among lowest ones for most quantiles τ . On the other hand, PLQF possesses one of the highest quantile scores across the same quantiles.

$$QS^k(\tau) = \frac{1}{T^*} \sum_{t=1}^{T^*} \left(r_{t+1} - \hat{Q}_{t+1,t}^k(\tau) \right) \left(1_{\{r_{t+1} \leq \hat{Q}_{t+1,t}^k(\tau)\}} - \tau \right) \quad (11)$$

where T^* is the number of out-of-sample forecasts, r_{t+1} is the realized value of equity premium, $\hat{Q}_{t+1,t}^k(\tau)$ represents the quantile forecast at level τ of model k , and indicator function $1_{\{\cdot\}}$ equals 1 if $r_{t+1} \leq \hat{Q}_{t+1,t}^k(\tau)$; otherwise it equals 0. As a result, quantile scores are always negative. Thus the larger QS is, i.e. the closer it is to zero, the better.

Table IV shows the QS for each single-predictor quantile model ($\hat{r}_{i,t+1}(\tau)$), AL-LASSO ($\hat{r}_{t+1}(\tau)$) and PLQF ($f_{t+1,t}^\tau$), over the full out-of-sample period 1967:1–2013:12. We see that AL-LASSO does not perform well because its quantile scores are among the lowest for most quantiles τ . On the other hand, PLQF possesses one of the highest quantile scores across the same quantiles τ . Moreover, none of the single-predictor quantile forecasts consistently outperforms PLQF across τ . Since accurate quantile forecasts are essential to yield successful point forecasts in the second step, the success of the PLQC point forecast relative to other quantile combination-based models is explained by the fact that it averages the most accurate quantile forecasts of equity premium.

Figure 6 shows the cumulative squared forecast error of the HA minus the cumulative squared forecast errors of point forecasts obtained by combining quantile forecasts from PLQF, AL-LASSO and single-predictor quantile models.¹⁹ We additionally report what Meligkotsidou *et al.* (2014) called robust forecast combination (RFC₁), which is computed by averaging all the 15 point forecasts obtained from the single-predictor quantile forecasting models.

Figure 6 suggests that the point forecasts obtained from single-predictor quantile models and AL-LASSO are still unable to outperform HA consistently over time in terms of their cumulative

¹⁹ For the sake of brevity and without affecting our conclusions, we only use the first weighting scheme to compute these point forecasts. Each single-predictor quantile forecasting model generates one point forecast. Thus there will be 15 such point forecasts.

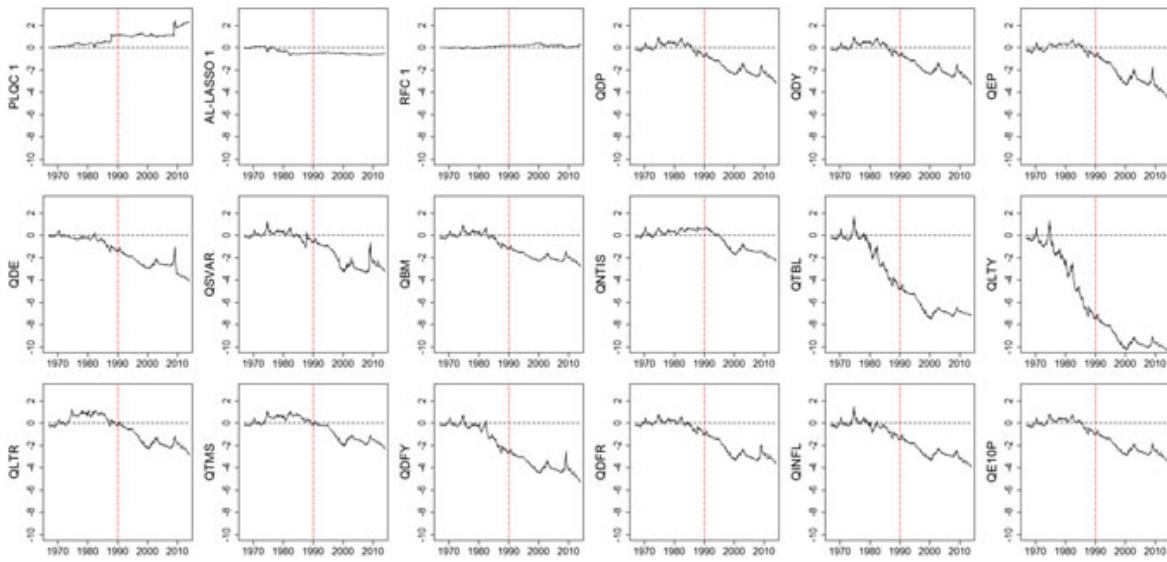


Figure 6. Cumulative squared prediction error for the benchmark model minus the cumulative squared prediction errors for the PLQC 1, AL-LASSO 1, RFC 1 and single-predictor quantile forecasting models, 1967:1–2013:12. A positively sloped curve in each panel indicates that the conditional model outperforms the HA, while the opposite holds for a downward-sloping curve. Moreover, if the curve is higher at the end of the period, the conditional model has a lower MSPE than the benchmark over this period. In this figure, the cumulative performance of single-predictor quantile models, AL-LASSO and RFC 1, hardly beat that of the HA consistently over time, as the PLQC forecast does. [Colour figure can be viewed at wileyonlinelibrary.com]

performance. RFC_1 hardly outperforms the historical average in any consistent basis of time. This happens because, unlike the PLQC forecast, these models are not designed to deal with partially and fully weak predictors across quantiles and over time, and thus are severely affected by misspecification. The failure of the AL-LASSO can also be explained by the fact that quantiles are not additive.²⁰. In other words, the AL-LASSO method assumes quantile additivity, $\hat{r}_{t+1}(\tau) = \sum_{j=1}^{15} \hat{\theta}_{i,t} \hat{r}_{i,t+1}(\tau)$, which may not hold in practice. The cumulative performance of PLQC forecast beats the HA over time and shows a clear superiority over other point (MSPE) forecasts obtained from a combination of quantile forecasts.

4. CONCLUSION

This paper studies equity premium forecasting using monthly observations of returns to the S&P 500 from 1926:12 to 2013:12. A common feature of existing models is that they produce inaccurate forecasts due to the presence of weak predictors and estimation errors. We propose a model selection procedure to identify partially and fully weak predictors, and use this information to make optimal MSPE forecasts based on an averaging scheme applied to quantiles. The quantiles combination, as a robust approximation to the conditional mean, avoids accuracy loss caused by estimation errors. The resulting PLQC forecasts achieve a middle ground in terms of variance versus bias, whereas existing methods reduce forecast variance significantly but are unable to lower bias by a large scale. For this reason, the PLQC forecast outperforms the historical average and other existing forecasting models by statistically and economically meaningful margins.

In the robustness analysis, we consider other quantile forecasting models based on fixed predictors. These models are not designed to deal with partial and fully weak predictors across quantiles and

²⁰ This means that for two random variables, X and Y , $Q_\tau(X + Y)$ is not necessarily equal to $Q_\tau(X) + Q_\tau(Y)$.

over time. The empirical results show that the quantile forecasts from such models are outperformed by the proposed post-LASSO quantile forecast (PLQF). Moreover, the point forecasts obtained from the combination of such quantile forecasts are still unable to provide a solution to the original puzzle reported by Welch and Goyal (2008).

In conclusion, equity premium forecasts can be improved if a method minimizes the effect of misspecification caused by weak predictors and estimation errors. Our results support the conclusion that an optimal MSPE out-of-sample forecast of the equity premium can be achieved when we integrate LASSO estimation and quantile combination into the same framework.

ACKNOWLEDGEMENTS

We thank the editor Jonathan Wright and two anonymous referees for their valuable and insightful comments. We also appreciate feedback from seminar participants at the Midwest Economic Association, the International Symposium on Forecasting, the University of Tennessee and the University of Illinois at Urbana-Champaign. We thank Asa Lambert and John McMahan for their careful reading of the paper.

REFERENCES

- Belloni A, Chernozhukov V. 2011. ℓ_1 -penalized quantile regression in high-dimensional sparse models. *Annals of Statistics* **39**: 82–130.
- Campbell JY. 2000. Asset pricing at the millennium. *Journal of Finance* **55**: 1515–1567.
- Campbell JY, Thompson SB. 2008. Predicting excess stock returns out of sample: can anything beat the historical average? *Review of Financial Studies* **21**: 1509–1531.
- Christoffersen PF, Diebold FX. 1997. Optimal prediction under asymmetric loss. *Econometric Theory* **13**(6): 808–817.
- Clark TE, McCracken MW. 2001. Tests of equal forecast accuracy and encompassing for nested models. *Journal of Econometrics* **105**(1): 85–110.
- Clark TE, West KD. 2007. Approximately normal tests for equal predictive accuracy in nested models. *Journal of Econometrics* **138**: 291–311.
- Diebold FX, Mariano RS. 1995. Comparing predictive accuracy. *Journal of Business and Economic Statistics* **13**: 253–263.
- Dow CH. 1920. Scientific stock speculation. *Magazine of Wall Street*.
- Elliott G, Gargano A, Timmermann A. 2013. Complete subset regressions. *Journal of Econometrics* **177**(2): 357–373.
- Gaglianone WP, Lima LR, Linton O, Smith DR. 2011. Evaluating value-at-risk models via quantile regression. *Journal of Business & Economic Statistics* **29**(1): 150–160.
- Gaglianone W, Lima LR. 2012. Constructing density forecasts from quantile regression. *Journal of Money, Credit and Banking* **44**: 1589–1607.
- Gaglianone WP, Lima LR. 2014. Constructing optimal density forecasts from point forecast combinations. *Journal of Applied Econometrics* **29**(5): 736–757.
- Granger CW. 1969. Investigating causal relations by econometric models and cross-spectral methods. *Econometrica* **37**(3): 424–438.
- Granger CW, Newbold P. 1986. *Forecasting Time Series* (2nd edn). Academic Press: New York.
- Issler JV, Lima LR. 2009. A panel data approach to economic forecasting: The bias-corrected average forecast. *Journal of Econometrics* **152**(2): 153–164.
- Judge GG, Hill RC, Griffiths WE, Lutkepohl H, Lee TC. 1988. *Introduction to the Theory and Practice of Econometrics*. Wiley: New York.
- Koenker R. 2005. *Quantile Regression*, Econometric Society Monographs No. 38. Cambridge University Press: Cambridge, UK.
- Koenker R, Machado JA. 1999. Goodness of fit and related inference processes for quantile regression. *Journal of the American Statistical Association* **94**(448): 1296–1310.
- Ma L, Pohlman L. 2008. Return forecasts and optimal portfolio construction: A quantile regression approach. *European Journal of Finance* **14**(5): 409–425.

- Manzan S. 2015. Forecasting the distribution of economic variables in a data-rich environment. *Journal of Business and Economic Statistics* **33**(1): 144–164.
- McCracken MW. 2007. Asymptotics for out of sample tests of Granger causality. *Journal of Econometrics* **140**(2): 719–752.
- Meligkotsidou L, Panopoulou E, Vrontos I, Vrontos S. 2014. A quantile regression approach to equity premium prediction. *Journal of Forecasting* **33**(7): 558–576.
- Patton AJ, Timmermann A. 2007a. Properties of optimal forecasts under asymmetric loss and nonlinearity. *Journal of Econometrics* **140**(2): 884–918.
- Patton AJ, Timmermann A. 2007b. Testing forecast optimality under unknown loss. *Journal of the American Statistical Association* **102**(480): 1172–1184.
- Pesaran MH, Timmermann A. 1995. Predictability of stock returns: Robustness and economic significance. *Journal of Finance* **50**(4): 1201–1228.
- Rapach D, Zhou G. 2013. Forecasting stock returns. In *Handbook of Economic Forecasting*, Timmermann A, Elliott G (eds), Part A. Elsevier: Amsterdam; 328–383.
- Rapach DE, Strauss JK, Zhou G. 2010. Out-of-sample equity premium prediction: Combination forecasts and links to the real economy. *Review of Financial Studies* **23**(2): 821–862.
- Taylor JW. 2007. Forecasting daily supermarket sales using exponentially weighted quantile regression. *European Journal of Operational Research* **178**: 154–167.
- Van de Geer. 2008. High-dimensional generalized linear models and the lasso. *Annals of Statistics* **36**(2): 614–645.
- Welch I, Goyal A. 2008. A comprehensive look at the empirical performance of equity premium prediction. *Review of Financial Studies* **21**: 1455–1508.
- West KD. 1996. Asymptotic inference about predictive ability. *Econometrica* **64**: 1067–1084.